

Java / J2EE Backend Senior Professional

Complete Interview Preparation Series

19 Chapters · ~965 Pages · Deep Q&A; + Code Examples + Exercises

Phase 1: Foundations

Phase 2: Framework & Data

Phase 3: Security

Phase 4: Capabilities

Phase 5: Perf & Architecture

Phase 6: Infrastructure

Targeting: Google · Amazon · Meta · Microsoft · Top-Tier Tech Companies

For Professionals with 15 – 20+ Years of Experience

**19
Chapters**

~965 Pages

6 Phases

**4 Exercise
Types**

Table of Contents

Phase 1: Foundations

- Chapter 01 — Core Java & JVM Internals (75 pages)
- Chapter 02 — Design Patterns & Design Principles (55 pages)
- Chapter 03 — Data Structures & Algorithms in Java (40 pages)
- Chapter 04 — Build Tools: Maven & Gradle (40 pages)

Phase 2: Application Framework & Data Layer

- Chapter 05 — Spring Framework: Internals-First (55 pages)
- Chapter 06 — JPA & ORM Internals (55 pages)
- Chapter 07 — Databases & SQL (40 pages)

Phase 3: Application Security

- Chapter 08 — Security: Auth, Certificates, SSO & API Security (45 pages)

Phase 4: Application Capabilities

- Chapter 09 — Caching: All Levels (50 pages)
- Chapter 10 — Application Search: Elasticsearch & Alternatives (45 pages)
- Chapter 11 — Internationalisation: Multi-Language, Multi-Currency & Multi-Timezone (45 pages)
- Chapter 12 — Concurrent User Edits: Conflict Detection & Graceful Handling (45 pages)

Phase 5: Performance & Architecture

- Chapter 13 — Software Performance: Sync, Async & Semi-Sync (45 pages)
- Chapter 14 — Microservices & Distributed Architecture (45 pages)
- Chapter 15 — System Design: Principles, Patterns & Real-World Problems (70 pages)

Phase 6: Infrastructure & Deployment

- Chapter 16 — Docker: Images, Containers & Docker Without Kubernetes (50 pages)
- Chapter 17 — Reverse Proxy: Nginx, Ingress & Traffic Management (50 pages)
- Chapter 18 — Kubernetes: Everything End-to-End (65 pages)
- Chapter 19 — DevOps, CI/CD, AWS & GCP (65 pages)

Total: 19 chapters · 980 pages

About This Series

This 19-chapter series is a complete interview preparation guide for senior Java / J2EE backend professionals targeting principal and staff engineer roles at top-tier technology companies such as Google, Amazon, Meta, and Microsoft. The series is structured in the exact order you encounter these topics when building real production software — from core language foundations through application frameworks, security, capabilities, performance, architecture, and finally infrastructure and cloud deployment.

Each chapter combines three elements: (1) concise documentation explaining how the technology works internally, (2) deep interview Q&A; with full answers and Java code examples, and (3) four types of exercises — coding challenges, debug/fix problems, design problems, and MCQ / scenario quizzes. Answers are targeted at 15-20+ years of experience and cover both senior and principal/staff level depth.

The series is delivered as 19 individual PDFs — one per chapter — so you can focus on the areas most relevant to your interview preparation without being overwhelmed by a single document.

Exercise Types Included in Every Chapter

Coding Challenges**Debug / Fix Problems****Design Problems****MCQ / Scenario Quiz**

Phase 1: Foundations

4 chapters · ~210 pages

01

Core Java & JVM Internals

Java fundamentals: primitives vs wrappers, autoboxing, pass-by-value, immutability, String pool, StringBuilder vs StringBuffer, equals()/hashCode() contract. Collections Framework (in-depth): Full interface/class hierarchy — Iterable > Collection > List/Set/Queue/Deque; Map hierarchy (separate). LIST: ArrayList (dynamic array, 1.5x growth, O(1) get), LinkedList (doubly-linked + Deque), CopyOnWriteArrayList (thread-safe snapshot); imports, all List/ArrayList methods (add, get, set, remove, indexOf, subList, sort, toArray, List.of), complete code examples. SET: HashSet (hash table, O(1)), LinkedHashSet (insertion order), TreeSet (red-black tree, sorted, NavigableSet — floor/ceiling/headSet/tailSet), EnumSet (bit-vector); imports, all Set/SortedSet/NavigableSet methods, set operations (union/intersection/difference), code examples. MAP: HashMap (bucket array, treeification at bin-size 8, ... (full detail in the chapter PDF)

Est. 75 pages

02

Design Patterns & Design Principles

Design Principles (OOP & SOLID): Single Responsibility (one reason to change), Open/Closed (open for extension, closed for modification), Liskov Substitution (subtypes must be substitutable — covariant return types, no precondition strengthening), Interface Segregation (many specific interfaces > one fat interface), Dependency Inversion (depend on abstractions not concretions). Additional principles: DRY (Don't Repeat Yourself), KISS (Keep It Simple), YAGNI (You Aren't Gonna Need It), SoC (Separation of Concerns), LoD / Law of Demeter (talk only to immediate friends), Composition over Inheritance, Program to Interfaces. Cohesion vs Coupling — measuring and improving each. Creational Patterns: Singleton (double-checked locking with volatile, enum singleton, holder idiom — thread safety analysis for each), Factory Method (decouple object creation, subclass decides), Abstract Factory ... (full detail in the chapter PDF)

Est. 55 pages

03

Data Structures & Algorithms in Java

Covers Big-O complexity analysis with a Java collection operation reference table, arrays and strings (sliding window, two-pointer, prefix sum), linked lists, stacks and queues, trees and BSTs (traversals, AVL), heaps and priority queues (k-way merge, top-K), hash tables (Java HashMap as implementation), graphs (BFS, DFS, Dijkstra, topological sort, union-find), sorting algorithms (merge, quick, heap sort with Java implementations), dynamic programming (memoisation vs tabulation, classic patterns), recursion and backtracking, and segment/Fenwick trees for range queries.

Est. 40 pages

04

Build Tools: Maven & Gradle

Covers Maven's POM, GAV coordinates, lifecycle phases, plugin system, dependency scopes, transitive dependency resolution, nearest-wins conflict rule, dependencyManagement, BOMs, and multi-module Reactor builds. Covers Gradle's task DAG, incremental builds, build cache, Kotlin DSL, configurations (implementation vs api), version catalogs, dependency locking, and multi-project builds. Includes a full Maven vs Gradle comparison (performance, flexibility, learning curve, CI integration), fat JAR packaging, publishing to Artifactory/Nexus, and build performance optimisation for CI pipelines.

Est. 40 pages

Phase 2: Application Framework & Data Layer

3 chapters · ~150 pages

05

Spring Framework: Internals-First

Covers IoC and DI fundamentals (BeanFactory vs ApplicationContext), all three configuration styles (XML, annotation, Java), @Autowired resolution internals via AutowiredAnnotationBeanPostProcessor, bean scopes and the prototype-in-singleton problem, the full Spring bean lifecycle (all 12 steps of ApplicationContext.refresh()), Spring proxy mechanism (JDK dynamic proxy vs CGLIB - when each is used, self-invocation pitfall), Spring AOP (advice types, pointcut expressions, ordering), @Transactional internals (proxy-based, all propagation and isolation levels), Spring MVC request pipeline (DispatcherServlet through ViewResolver), Spring Boot auto-configuration (@Conditional annotations, custom starters), Spring WebFlux reactive model (Mono/Flux, backpressure), and Spring Security internals (filter chain, OAuth2, method-level security).

Est. 55 pages

06

JPA & ORM Internals

Covers JPA spec vs Hibernate vs Spring Data JPA, EntityManagerFactory and EntityManager internals, persistence context as identity map, all four entity lifecycle states and transitions, dirty checking and flush modes, first-level cache (session cache) and batch processing implications, second-level cache (providers, concurrency strategies), all association types with fetch strategies, the N+1 problem and every solution (JOIN FETCH, @EntityGraph, @BatchSize, DTO projections), inheritance mapping strategies (SINGLE_TABLE, JOINED, TABLE_PER_CLASS) with trade-offs, optimistic and pessimistic locking with @Version, bulk operations and batch insert performance, JPQL vs Criteria API vs native queries, advanced Hibernate features (@Formula, @Where, @Filter, @NaturalId, JSON columns), and Hibernate 6 / Jakarta Persistence 3 changes.

Est. 55 pages

07

Databases & SQL

Covers JDBC (Connection, PreparedStatement, connection pooling), SQL deep dive (joins, subqueries, CTEs, window functions), B-tree vs LSM tree index structures, clustered vs non-clustered indexes, covering indexes and composite index prefix matching, query execution plan analysis (EXPLAIN ANALYZE), ACID properties and how databases guarantee them, all four isolation levels and what anomalies each prevents, MVCC (how readers don't block writers), deadlock detection and prevention, database design (normalisation, when to denormalise), NoSQL databases (Cassandra, MongoDB, DynamoDB, Redis), schema migration with Flyway and Liquibase, and read replica and sharding strategies.

Est. 40 pages**Phase 3: Application Security**

1 chapters · ~45 pages

08

Security: Auth, Certificates, SSO & API Security

Covers cryptography fundamentals (AES, RSA, hashing, digital signatures, HMAC), HTTPS and TLS internals (TLS 1.3 handshake step by step, cipher suites, forward secrecy), X.509 certificates and PKI (certificate chain, CA trust model, CRL vs OCSP, KeyStore/TrustStore in Java, mTLS), cookies in depth (HttpOnly, Secure, SameSite, CSRF relationship), authentication fundamentals (password storage with bcrypt, MFA/TOTP, Spring Security), authorisation models (RBAC, ABAC, ReBAC), OAuth2 deep dive (all grant types, PKCE), OpenID Connect (ID token, discovery endpoint), JWT (structure, RS256 vs HS256, alg:none vulnerability, revocation), Single Sign-On (SAML 2.0, OIDC-based SSO, session federation, Single Logout, Keycloak, LDAP/Active Directory), OWASP Top 10 with Java-specific mitigations, and API security best practices.

Est. 45 pages**Phase 4: Application Capabilities**

4 chapters · ~185 pages

09

Caching: All Levels

Covers the memory hierarchy and latency numbers, L1/L2/L3 CPU cache and false sharing (@Contended in Java), in-process JVM cache with Caffeine (W-TinyLFU eviction algorithm), distributed cache with Redis (data structures, persistence, Cluster, Sentinel, Pub/Sub), Hibernate first-level and second-level cache, database query cache and materialised views, CDN and edge caching (Cache-Control, ETag, surrogate keys, stale-while-revalidate), HTTP and browser cache, reverse proxy cache (Nginx, Varnish, API Gateway caching), multi-level tiered caching architecture (L1 Caffeine + L2 Redis + L3 DB), all five caching strategies (cache-aside, read-through, write-through, write-behind, refresh-ahead), cache invalidation approaches, all eviction policies (LRU, LFU, W-TinyLFU, TTL), and Spring @Cacheable/@CacheEvict integration.

Est. 50 pages

10

Application Search: Elasticsearch & Alternatives

Covers why database LIKE queries fail at scale, inverted index internals (tokenisation pipeline, stemming, synonyms, posting lists), Elasticsearch architecture (cluster, shards, segments, write/read path, NRT indexing, ILM hot/warm/cold tiers), Query DSL in depth (term vs full-text queries, bool query, geo queries, function score), relevance scoring (BM25, field boosting, explain API), Spring Data Elasticsearch integration, keeping Elasticsearch in sync with the database (dual-write, Debezium CDC + Kafka, zero-downtime reindex with aliases), advanced features (autocomplete, fuzzy search, vector/semantic search with kNN), alternative search engines (Solr, OpenSearch, Algolia, Typesense, PostgreSQL tsvector, pgvector), and search architecture patterns (CQRS applied to search, search analytics, personalised ranking).

Est. 45 pages

11

Internationalisation: Multi-Language, Multi-Currency & Multi-Timezone

Covers i18n vs L10n vs Globalisation, Java Locale and Accept-Language header handling, Spring MessageSource and LocaleResolver, resource bundles with pluralisation (ICU MessageFormat), database content localisation strategies, text encoding (UTF-8 everywhere, utf8mb4 in MySQL), bidirectional text (RTL for Arabic/Hebrew, CSS direction, Unicode BiDi algorithm), BigDecimal for money (never double), JSR 354 Moneta library, ISO 4217 currency codes, exchange rate storage and caching, locale-aware currency formatting, rounding modes, the golden rule of timezone handling (store UTC, display local), java.time API in depth (Instant vs ZonedDateTime vs LocalDateTime), DST pitfalls and recurring event scheduling, user timezone detection, Jackson timezone serialisation, NumberFormat and DateTimeFormatter with locale, and Elasticsearch language analysers.

Est. 45 pages

12

Concurrent User Edits: Conflict Detection & Graceful Handling

Covers the lost update problem and silent overwrite (the default broken behaviour), all six conflict handling strategies: Last-Write-Wins (when acceptable), Optimistic Concurrency Control with @Version (end-to-end flow from UI to DB), HTTP ETags and conditional requests (If-Match, 412 Precondition Failed), Pessimistic Locking (record checkout with heartbeat, lock expiry, admin override), Field-Level Merge with three-way diff (JSON Patch, @DynamicUpdate), and Queue-Based Serialisation (Kafka partition key per entity, actor model). Covers graceful error handling - converting OptimisticLockException to structured 409 Conflict with user-facing diff, Spring @ExceptionHandler with RFC 9457 ProblemDetail, UI retry-with-reload patterns, idempotency keys for safe retries, and full end-to-end scenario walkthroughs (hotel double-booking, CRM edit conflict, flash sale).

Est. 45 pages**Phase 5: Performance & Architecture**

3 chapters · ~160 pages

13

Software Performance: Sync, Async & Semi-Sync

Covers performance fundamentals (latency vs throughput, percentile metrics, Little's Law, Amdahl's Law), synchronous communication in depth (thread-per-request model, connection pool sizing, cascading latency, C10K problem - pros and cons), asynchronous communication (event-loop model, NIO, Project Reactor, CompletableFuture, async messaging - pros and cons), semi-synchronous patterns (Future with timeout, async-then-poll, request-reply over Kafka, virtual threads as 'sync code with async performance' - pros and cons), threading models compared (thread-per-request vs event loop vs virtual threads vs Fork/Join vs actor model), JVM performance tuning (GC selection, heap sizing, JIT warm-up, GraalVM native, profiling with JFR and flame graphs), database performance (connection pool exhaustion, N+1 detection, read replicas, write batching), and load/stress/soak/spike testing with Gatling ... (full detail in the chapter PDF)

Est. 45 pages

14

Microservices & Distributed Architecture

Covers microservice decomposition principles (DDD bounded contexts, single responsibility, right-sizing), all inter-service communication patterns (REST, gRPC, async messaging), distributed transactions without 2PC (Saga pattern - choreography vs orchestration, Outbox pattern), API Gateway responsibilities (TLS termination, auth offloading, rate limiting, routing, BFF pattern), service discovery, resilience patterns (Circuit Breaker with Resilience4j, Bulkhead thread pool isolation, Retry with exponential backoff and jitter, Rate Limiter token vs leaky bucket), observability (distributed tracing, structured logging, Micrometer + Prometheus + Grafana, SLI/SLO/SLA), event-driven architecture (event vs command vs query, schema evolution), Kafka deep dive (ISR, exactly-once semantics, consumer groups, KEDA scaling), distributed systems fundamentals (CAP/PACELC, Raft consensus, vector ... (full detail in the chapter PDF)

Est. 45 pages

15

System Design: Principles, Patterns & Real-World Problems

A comprehensive system design chapter covering every dimension expected at Google, Amazon, Meta, and Microsoft principal/staff engineer interviews. Foundations: How to approach a system design interview (clarify requirements, estimate scale, define APIs, design data model, high-level architecture, deep dive, identify bottlenecks). Back-of-the-envelope estimation — latency numbers every engineer must know (L1 cache 0.5ns, RAM 100ns, SSD 100us, network roundtrip 150ms), QPS estimation, storage calculation, bandwidth planning. Scalability Principles: Horizontal vs vertical scaling trade-offs, stateless vs stateful services (why statelessness enables horizontal scale), database connection pool sizing, the Fan-Out problem (social network feed distribution — push vs pull model trade-offs), back pressure as a concurrency control mechanism. Load Balancing: L4 (TCP) vs L7 (HTTP) load balancers, ... (full detail in the chapter PDF)

Est. 70 pages

Phase 6: Infrastructure & Deployment

4 chapters · ~230 pages

16

Docker: Images, Containers & Docker Without Kubernetes

Covers containers vs VMs (namespaces, cgroups, union filesystem), Docker architecture (daemon, CLI, containerd, runc, OCI standard), image layers and the overlay2 filesystem (content-addressed SHA256 layers, layer caching strategy), Dockerfile best practices for Java (multi-stage builds reducing image from 800MB to 120MB, exec form vs shell form for signal handling, non-root user, MaxRAMPercentage for container awareness), Docker networking (bridge, host, user-defined bridge with DNS, overlay), volumes (bind mount vs named volume vs tmpfs), Docker Compose for multi-container local development, image security (Trivy scanning, distroless/Chainguard images, image signing with Cosign), Docker in CI/CD (BuildKit, multi-platform builds with buildx), Docker Swarm for simple multi-server orchestration, and a comprehensive answer to 'Can Docker be effective without Kubernetes?' covering Docker ... (full detail in the chapter PDF)

Est. 50 pages

17

Reverse Proxy: Nginx, Ingress & Traffic Management

Covers forward proxy vs reverse proxy (what each faces, why the names), the full set of reverse proxy responsibilities (TLS termination, load balancing, routing, compression, caching, rate limiting, auth offloading, header transformation), Nginx architecture (master/worker, event-driven non-blocking I/O), nginx.conf structure (http/server/location blocks), location matching order, proxy_pass and upstream configuration, SSL setup, rate limiting with limit_req_zone, annotated Nginx config examples for Spring Boot microservices, Kubernetes Service types compared (ClusterIP, NodePort, LoadBalancer, ExternalName, Headless), the Nginx Ingress Controller architecture (how it dynamically generates nginx.conf), a complete annotated Kubernetes architecture showing internal-only services (postgres, redis, workers) vs externally-routed services (booking-api, admin-ui) via Ingress, full Ingress YAML ... (full detail in the chapter PDF)

Est. 50 pages

18

Kubernetes: Everything End-to-End

Covers why Kubernetes exists (desired state reconciliation), full control plane internals (kube-apiserver, etcd/Raft, scheduler, controller-manager, cloud-controller-manager), worker node components (kubelet, kube-proxy, containerd/runc), every workload type (Pod, ReplicaSet, Deployment, StatefulSet, DaemonSet, Job, CronJob, HPA, VPA) with use cases, all Service types, storage (PV, PVC, StorageClass, EBS/PD CSI drivers), ConfigMaps and Secrets (External Secrets Operator for AWS/GCP Secrets Manager), full deployment on AWS EKS (eksctl/Terraform, managed node groups vs Fargate, AWS Load Balancer Controller, IRSA, VPC CNI, Karpenter), full deployment on GCP GKE (Standard vs Autopilot, Workload Identity, GCP HTTP(S) LB Ingress, BackendConfig, NEG's, Cloud Armor), programmatically invoking the Kubernetes API for scaling (Fabric8 Java client - scale Deployments, patch HPA, watch rollout, KEDA ... (full detail in the chapter PDF)

Est. 65 pages

19

DevOps, CI/CD, AWS & GCP

Covers the 12-Factor App methodology applied to Spring Boot, immutable infrastructure, CI/CD pipeline stages (build, unit test, integration test, security scan, staging, production), trunk-based development vs GitFlow, deployment strategies (blue-green, canary, feature flags, rolling), safe database migrations in CD pipelines, rollback strategies, shift-left testing, AWS services for Java backends in depth (EC2/ECS/EKS/Lambda, S3, RDS/Aurora, DynamoDB, ElastiCache, SQS/SNS/EventBridge/Kinesis, API Gateway, CloudFront, Route 53, VPC/IAM/Cognito/Secrets Manager/KMS, CloudWatch/X-Ray, CodePipeline/CDK), GCP services for Java backends in depth (GKE/Cloud Run/Cloud Functions, Cloud SQL/AlloyDB/Spanner/Firestore/Bigtable/Memorystore, Pub/Sub/Cloud Tasks/Workflows, Cloud Load Balancing/CDN/Cloud Armor, Cloud IAM/Secret Manager, Cloud Build/Cloud Deploy/Artifact Registry, BigQuery/Dataflow), ... (full detail in the chapter PDF)

Est. 65 pages

Series at a Glance

#	Chapter	Phase	Pages
1	Core Java & JVM Internals	<i>Ph 1: Foundations</i>	~75
2	Design Patterns & Design Principles	<i>Ph 1: Foundations</i>	~55
3	Data Structures & Algorithms in Java	<i>Ph 1: Foundations</i>	~40
4	Build Tools: Maven & Gradle	<i>Ph 1: Foundations</i>	~40
5	Spring Framework: Internals-First	<i>Ph 2: Application Framework & Data Layer</i>	~55
6	JPA & ORM Internals	<i>Ph 2: Application Framework & Data Layer</i>	~55
7	Databases & SQL	<i>Ph 2: Application Framework & Data Layer</i>	~40
8	Security: Auth, Certificates, SSO & API Security	<i>Ph 3: Application Security</i>	~45
9	Caching: All Levels	<i>Ph 4: Application Capabilities</i>	~50
10	Application Search: Elasticsearch & Alternatives	<i>Ph 4: Application Capabilities</i>	~45
11	Internationalisation: Multi-Language, Multi-Currency & Multi-Timezone	<i>Ph 4: Application Capabilities</i>	~45
12	Concurrent User Edits: Conflict Detection & Graceful Handling	<i>Ph 4: Application Capabilities</i>	~45
13	Software Performance: Sync, Async & Semi-Sync	<i>Ph 5: Performance & Architecture</i>	~45
14	Microservices & Distributed Architecture	<i>Ph 5: Performance & Architecture</i>	~45
15	System Design: Principles, Patterns & Real-World Problems	<i>Ph 5: Performance & Architecture</i>	~70
16	Docker: Images, Containers & Docker Without Kubernetes	<i>Ph 6: Infrastructure & Deployment</i>	~50
17	Reverse Proxy: Nginx, Ingress & Traffic Management	<i>Ph 6: Infrastructure & Deployment</i>	~50
18	Kubernetes: Everything End-to-End	<i>Ph 6: Infrastructure & Deployment</i>	~65
19	DevOps, CI/CD, AWS & GCP	<i>Ph 6: Infrastructure & Deployment</i>	~65
	Total		~980